

Complications - I

Reality 1: Cost of remote data access often not uniform

- often takes the form $Ax + B$ for x bytes
- may be more complex as x gets larger (e.g. TCP file transfer)

Reality 2: “Nearby” objects often accessed soon after object access

- another empirical observation about real systems in real use
- referred to as “*spatial locality of reference*” or “*spatial locality*”

Reality 3: Remote data more coarsely addressable than local

- typically a scalability tradeoff at next level of memory hierarchy
same number of address bits can cover larger volume of data
e.g. cache line width, page size, whole file, tape mount
- “wholesale” versus “retail”

Combining observations \Rightarrow *fetch more than you need on miss*


- effectively amortizes cost of fetch
- assumption sometimes violated
- extra data fetched is then useless (may even hurt)
can be viewed as crude form of prefetching

Complications - II

Temporal and spatial locality are very different properties

- caching implementations often tightly combine these assumptions
- one can exist without the other
 - spatial without temporal: linear scan of huge file
 - temporal without spatial: tight loop accessing just one object

For example, consider typical implementation of “rm -f *”



shell expands “*” into list

loop iterates through list

- stat object
- unlink object

parent directory exhibits temporal locality

directory entries exhibit spatial locality

Complications - III

Local storage management

- overhead typically makes first reference more expensive
local copy allocation, lookup table update, etc.
- remote storage often much bigger than local storage
⇒ recycling of local storage for copies
- **replacement** policy becomes significant

Updates to remote copy need to be propagated to local copy

- and vice versa (local updates need to be made visible everywhere)
- **cache consistency** is a significant problem
- goal is **one-copy semantics**

Refinements of Basic Idea

Cache idea can be applied recursively \Rightarrow “*multi-level cache*”
or “*memory hierarchy*”

Persistent caches (typically on-disk) increase longevity

Local copy can be used to mask remote failures \Rightarrow
caching for disconnected operation

Cooperative caching (across users) exploits “*communal locality*”

- empirical observation about groups of users

Outsourced caching to third parties (e.g. Akamai)

Users notice long fetch delays \rightarrow *translucent caching*

A Brief History of Caching

Demand paging was first known use of caching idea (1961)

- *"Dynamic Storage Allocation in the Atlas Computer, Including an Automatic Use of a Backing Store,"* John Fotheringham, Communications of the ACM, 1961, pp 435-436
- revisited by OS researchers (1970's - 1990s)

Hardware caches came next (1968)

- *"Structural Aspects of the System/360 Model 85, Part II: The Cache,"* J. S. Liptay, IBM Systems Journal, Vol. 7, No. 1, 1968, pp. 15-21
- extensive study of hardware cache coherence (1970-1990s)

Distributed file systems came after that (~1983-84, ~1988-1991)

- *AFS, NFS, Sprite, Coda*
- extensive study by researchers (1980's - present)

Web caching (~1994)

- cooperative caching, outsourced caching (mid-1990s to present)

Virtual machine state caching (~2002-2008)

- *Internet Suspend/Resume*

Result caching in discard-based search (~2007-2009)

- *Diamond*