


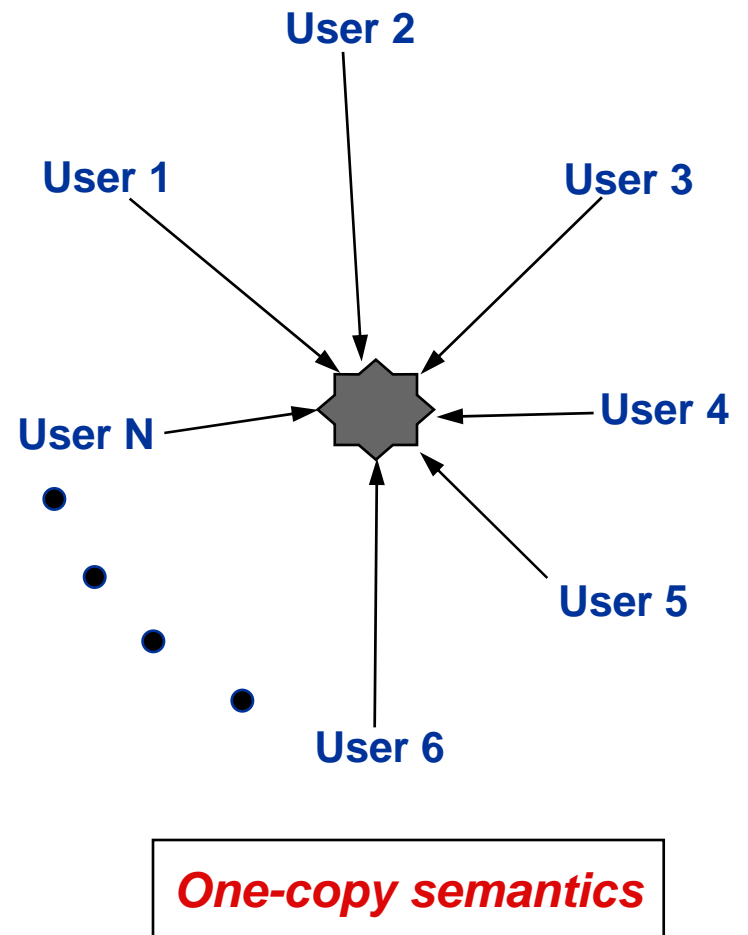
Update Propagation in a Cache-based System

(aka “Cache Consistency”)

Caching Reality



Desired Illusion



What Makes This Hard?

Physical master copy may not exist

- hosts track who has most recent copy
- more likely in P2P scenarios than client-server scenarios
- also common in multiprocessor hardware caches

Network may break between some users and master copy

- disconnected sites see no further updates
- other sites don't see updates by disconnected site

Intense read- and write-sharing across sites

- generates huge amount of cache propagation traffic
- interconnect becomes bottleneck for cache access
- neither writer-bias (write-back) nor reader-bias (write-through) helps
- caches effectively useless

Required Readings

- **[Howard88]**
Howard, J. H., Kazar, M. L., Menees, S. G., Nichols, D. A., Satyanarayanan, M., Sidebotham, R. N., and West, M. J. 1988. Scale and performance in a distributed file system. *ACM Trans. Comput. Syst.* 6, 1 (Feb. 1988)
- **[Gray89]**
Gray, C. and Cheriton, D. 1989. Leases: an efficient fault-tolerant mechanism for distributed file cache consistency. In *Proceedings of the Twelfth ACM Symposium on Operating Systems Principles*, Litchfield Park, AZ

Optional Readings

- **[Strenstrom90]**
P. Stentrsum, "A Survey of Cache Coherence Schemes for Multiprocessors", IEEE Computer, June 1990
- **[Satya85]**
Satyanarayanan, M., Howard, J. H., Nichols, D. A., Sidebotham, R. N., Spector, A. Z., and West, M. J. 1985. The ITC distributed file system: principles and design. In Proceedings of the Tenth ACM Symposium on Operating Systems Principles, Orcas Island, WA
- **[Nelson88]**
Nelson, M. N., Welch, B. B., and Ousterhout, J. K. 1988. Caching in the Sprite network file system. ACM Trans. Comput. Syst. 6, 1 (Feb. 1988)
- **[Luotonen94]**
Luotonen, A. and Altis, K. 1994. World-Wide Web proxies. In Selected Papers of the First Conference on World-Wide Web (Geneva, Switzerland). R. Cailliau and P. H. Enslow, Eds. Elsevier Science Publishers B. V., Amsterdam, The Netherlands
- **[Franklin97]**
Franklin, M. J., Carey, M. J., and Livny, M. 1997. Transactional client-server cache consistency: alternatives and performance. ACM Trans. Database Syst. 22, 3 (Sep. 1997)
- **[Yin02]**
Yin, J., Alvisi, L., Dahlin, M., and Iyengar, A. 2002. Engineering web cache consistency. ACM Trans. Inter. Tech. 2, 3 (Aug. 2002),

1. Broadcast Invalidations

Basic Idea

Every potential caching site notified on every update

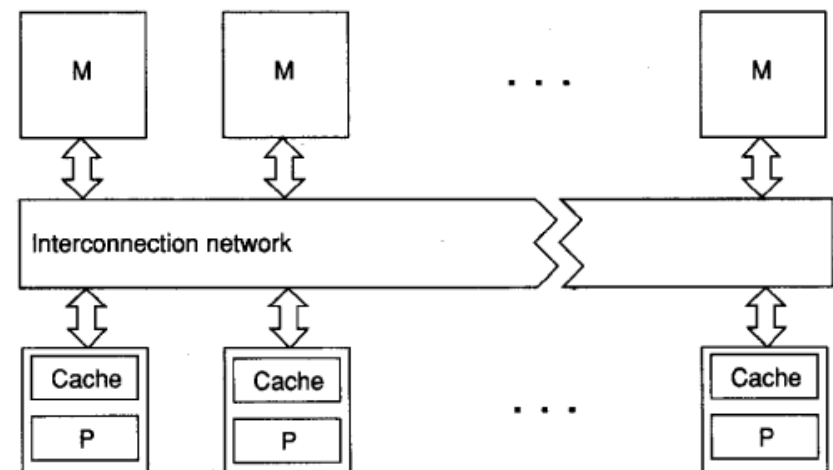
- No check to verify caching site actually contains object
- Notification includes specific object being invalidated
- Effectively broadcast of address being modified

At each cache site, next reference to object will cause a miss

Original Use

- Small-scale multiprocessors in early 1970's
- “snoopy cache” multiprocessor variants of 1980s-1990s

Reference: *Stenstrom90*



Strengths

- **very strict emulation of “one-copy semantics”**
- **no race conditions if updater blocked until all caches invalidated**
- **simple to implement**

Limitations

- **wasted traffic if no readers elsewhere**
- **updating process blocked until invalidation complete**
- **not a scalable design**

2. Check on Use

Basic Idea

- *reader checks master copy before each use*
- conditional fetch, if cache copy stale
- typically done at coarse granularity (e.g. entire file)

Original Use

- AFS-1 (circa 1983)
- Reference: *Satya85*

Advantages

- Strict consistency at coarse granularity
- Easy to implement, no server state
- Servers don't need to know of caching sites

Disadvantages

- Slows read access on loaded servers & high-latency networks
- Check is almost always success → frivolous traffic
- Load on network and server

3. Callback

Basic Idea


- *targeted notification of caching sites*
- master copy tracks sites with cached data
- typically done at coarse granularity (e.g. entire file)
- on update, all sites with cached data notified (*“callback”*)

Original Use

- AFS-2 (circa 1985)
- Essentially the same idea in MP systems → *“Directory scheme”*
- Reference: *Howard88, Stenstrom90*

Advantages

- Excellent scalability for Unix workloads
- *Zero network traffic for read of cached-valid objects*
- Precursor to caching for disconnected operation
- Biases read performance in favor of write-performance



Disadvantages

- sizable state on server
- complexity of tracking cached state on clients
- silence ambiguous for client: network failure → lost callbacks
 - periodic “keepalive” probes
 - data could be stale between probes
- NAT networks with masquerading firewalls